

```
#import library
from ast import increment_lineno
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', None)

from sklearn.feature_selection import RFE
from sklearn.linear_model import LinearRegression
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

from google.colab import files
uploaded=files.upload()
```

No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving CarPrice Assignment.csv to CarPrice Assignment.csv

```
car_dataset=pd.read_csv('CarPrice_Accommodation.csv',encoding='latin-1')
print('car_dataset')
```

```
car_dataset
```

```
car_dataset.head()
```

| | car_ID | symboling | CompanyName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation |
|---|--------|-----------|-------------|----------|------------|------------|-------------|------------|----------------|
| 0 | 1 | 3 | alfa-romero | gas | std | two | convertible | rwd | fr |
| 1 | 2 | 3 | alfa-romero | gas | std | two | convertible | rwd | fr |
| 2 | 3 | 1 | alfa-romero | gas | std | two | hatchback | rwd | fr |
| 3 | 4 | 2 | audi | gas | std | four | sedan | fwd | fr |
| 4 | 5 | 2 | audi | gas | std | four | sedan | 4wd | fr |

```
car_dataset.tail()
```

| | car_ID | symboling | CompanyName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation |
|-----|--------|-----------|-------------|----------|------------|------------|---------|------------|----------------|
| 200 | 201 | -1 | volvo | gas | std | four | sedan | rwd | fi |
| 201 | 202 | -1 | volvo | gas | turbo | four | sedan | rwd | fi |
| 202 | 203 | -1 | volvo | gas | std | four | sedan | rwd | fi |
| 203 | 204 | -1 | volvo | diesel | turbo | four | sedan | rwd | fi |
| 204 | 205 | -1 | volvo | gas | turbo | four | sedan | rwd | fi |

```
car_dataset.shape
```

```
(205, 26)
```

```
car_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   car_ID          205 non-null    int64  
 1   symboling       205 non-null    int64  
 2   CarName         205 non-null    object 

```

```

3   fueltype          205 non-null    object
4   aspiration        205 non-null    object
5   doornumber        205 non-null    object
6   carbody           205 non-null    object
7   drivewheel        205 non-null    object
8   enginolocation    205 non-null    object
9   wheelbase          205 non-null    float64
10  carlength         205 non-null    float64
11  carwidth          205 non-null    float64
12  carheight         205 non-null    float64
13  curbweight        205 non-null    int64
14  enginetype        205 non-null    object
15  cylindernumber    205 non-null    object
16  enginesize         205 non-null    int64
17  fuelsystem         205 non-null    object
18  boreratio          205 non-null    float64
19  stroke             205 non-null    float64
20  compressionratio   205 non-null    float64
21  horsepower         205 non-null    int64
22  peakrpm            205 non-null    int64
23  citympg            205 non-null    int64
24  highwaympg         205 non-null    int64
25  price              205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```
car_dataset.describe()
```

| | car_ID | symboling | wheelbase | carlength | carwidth | carheight | curbweight | enginesize | b |
|-------|------------|------------|------------|------------|------------|------------|-------------|------------|----|
| count | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 205.000000 | 20 |
| mean | 103.000000 | 0.834146 | 98.756585 | 174.049268 | 65.907805 | 53.724878 | 2555.565854 | 126.907317 | |
| std | 59.322565 | 1.245307 | 6.021776 | 12.337289 | 2.145204 | 2.443522 | 520.680204 | 41.642693 | |
| min | 1.000000 | -2.000000 | 86.600000 | 141.100000 | 60.300000 | 47.800000 | 1488.000000 | 61.000000 | |
| 25% | 52.000000 | 0.000000 | 94.500000 | 166.300000 | 64.100000 | 52.000000 | 2145.000000 | 97.000000 | |
| 50% | 103.000000 | 1.000000 | 97.000000 | 173.200000 | 65.500000 | 54.100000 | 2414.000000 | 120.000000 | |
| 75% | 154.000000 | 2.000000 | 102.400000 | 183.100000 | 66.900000 | 55.500000 | 2935.000000 | 141.000000 | |
| max | 205.000000 | 3.000000 | 120.900000 | 208.100000 | 72.300000 | 59.800000 | 4066.000000 | 326.000000 | |

```
car_dataset.isnull().sum()
```

```

car_ID          0
symboling       0
CarName         0
fueltype        0
aspiration      0
doornumber      0
carbody         0
drivewheel      0
enginolocation  0
wheelbase        0
carlength        0
carwidth         0
carheight        0
curbweight       0
enginetype       0
cylindernumber  0
enginesize       0
fuelsystem       0
boreratio        0
stroke           0
compressionratio 0
horsepower       0
peakrpm          0
citympg          0
highwaympg       0
price            0
dtype: int64

```

```
car_dataset.head(5)
```

| car_ID | symboling | CompanyName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation |
|--------|-----------|-------------|-------------|------------|------------|---------|-------------|----------------|
| 0 | 1 | 3 | alfa-romero | gas | std | two | convertible | rwd |
| 1 | 2 | 3 | alfa-romero | gas | std | two | convertible | rwd |
| 2 | 3 | 1 | alfa-romero | gas | std | two | hatchback | rwd |
| 3 | 4 | 2 | audi | gas | std | four | sedan | fwd |
| 4 | 5 | 2 | audi | gas | std | four | sedan | 4wd |

```
car_dataset.columns
```

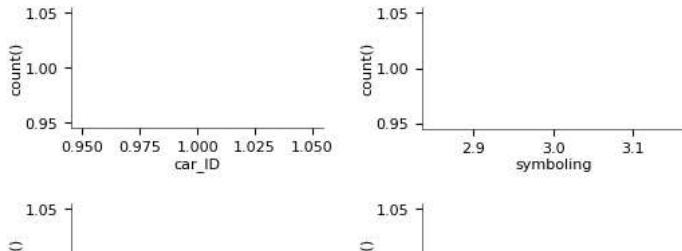
```
Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
       'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
       'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
       'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
       'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
       'price'],
      dtype='object')
```

```
CompanyName = car_dataset['CarName'].apply(lambda x : x.split(' ')[0])
car_dataset.insert(3,"CompanyName",CompanyName)
car_dataset.drop(['CarName'],axis=1,inplace=True)
```

```
car_dataset.head(1)
```

| car_ID | symboling | CompanyName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation |
|--------|-----------|-------------|-------------|------------|------------|---------|-------------|----------------|
| 0 | 1 | 3 | alfa-romero | gas | std | two | convertible | rwd |

Time series



```
car_dataset.CompanyName.unique()
```

```
array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
       'isuzu', 'jaguar', 'maxda', 'mazda', 'buick', 'mercury',
       'mitsubishi', 'nissan', 'nissan', 'peugeot', 'plymouth', 'porsche',
       'porcshce', 'renault', 'saab', 'subaru', 'toyota', 'toyouta',
       'vokswagen', 'volkswagen', 'vw', 'volvo'], dtype=object)
```

```
car_dataset.CompanyName = car_dataset.CompanyName.str.lower()
```

```
car_dataset.CompanyName.unique()
```

```
array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
       'isuzu', 'jaguar', 'maxda', 'mazda', 'buick', 'mercury',
       'mitsubishi', 'nissan', 'peugeot', 'plymouth', 'porsche',
       'porcshce', 'renault', 'saab', 'subaru', 'toyota', 'toyouta',
       'vokswagen', 'volkswagen', 'vw', 'volvo'], dtype=object)
```

```
def replace_name(a,b):
    car_dataset.CompanyName.replace(a,b,inplace=True)
```

```
replace_name('maxda','mazda')
replace_name('porcshce','porsche')
replace_name('toyouta','toyota')
replace_name('vokswagen','volkswagen')
replace_name('vw','volkswagen')
```

```
car_dataset.CompanyName.unique()
```

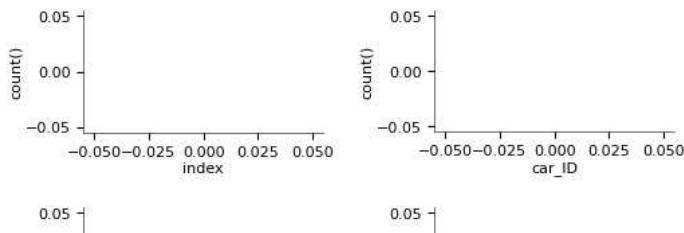
```
array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
       'isuzu', 'jaguar', 'mazda', 'buick', 'mercury', 'mitsubishi',
       'nissan', 'peugeot', 'plymouth', 'porsche', 'renault', 'saab',
       'subaru', 'toyota', 'volkswagen', 'volvo'], dtype=object)
```

```
car_dataset.loc[car_dataset.duplicated()]
```

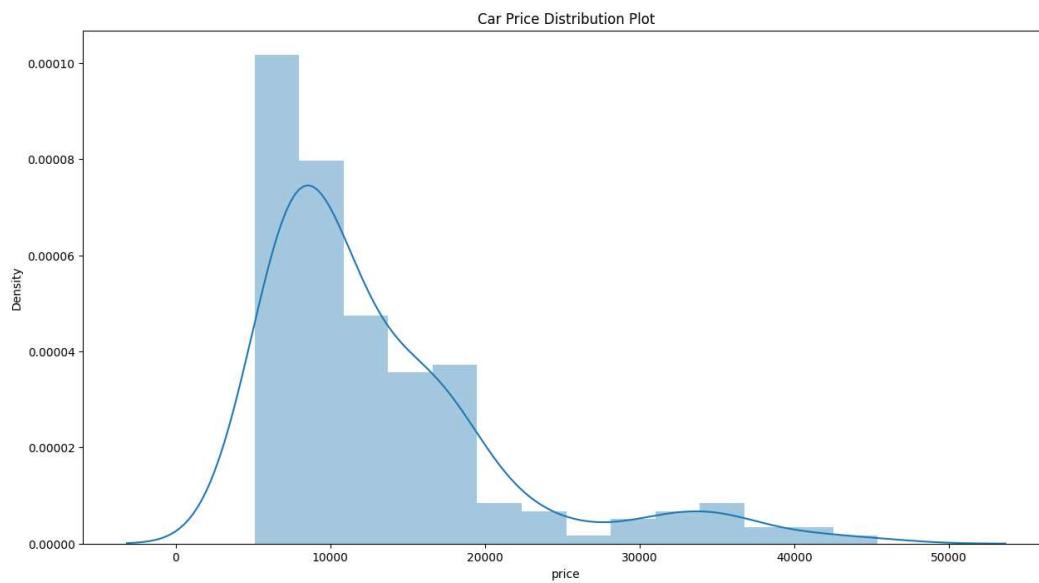
| car_ID | symboling | CompanyName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation |
|--------|-----------|-------------|----------|------------|------------|---------|------------|----------------|
|--------|-----------|-------------|----------|------------|------------|---------|------------|----------------|

Warning: Total number of columns (26) exceeds max_columns (20) limiting to first (20) columns.

Time series



```
plt.figure(figsize=(15,8))
plt.title('Car Price Distribution Plot')
sns.distplot(car_dataset.price)
plt.show()
```



```
car_dataset.price.describe(percentiles = [0.25,0.50,0.75,0.85,0.90,1])
```

| | |
|-------|--------------|
| count | 205.000000 |
| mean | 13276.710571 |
| std | 7988.852332 |
| min | 5118.000000 |
| 25% | 7788.000000 |

```

50%      10295.000000
75%      16503.000000
85%      18500.000000
90%      22563.000000
100%     45400.000000
max      45400.000000
Name: price, dtype: float64

```

```

plt.figure(figsize=(25, 8))

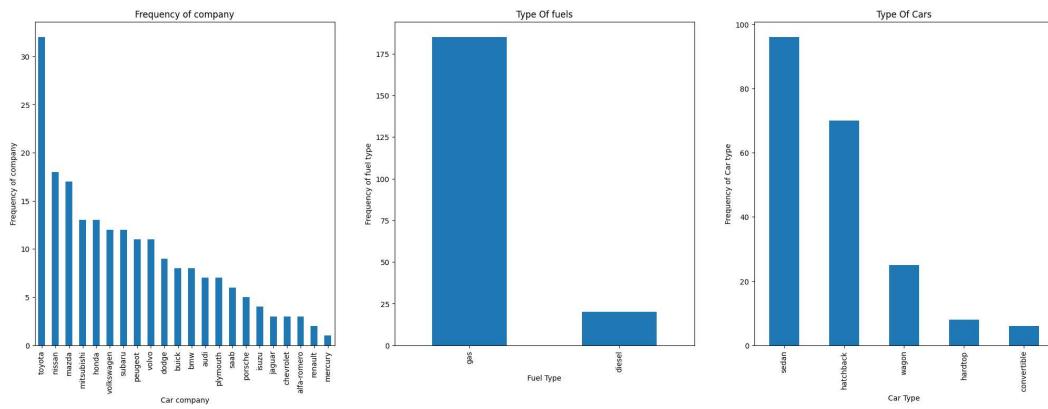
plt.subplot(1,3,1)
plt1 = car_dataset.CompanyName.value_counts().plot(kind='bar')
plt.title('Frequency of company')
plt1.set(xlabel = 'Car company', ylabel='Frequency of company')

plt.subplot(1,3,2)
plt1 = car_dataset.fueltype.value_counts().plot(kind='bar')
plt.title('Type Of fuels')
plt1.set(xlabel = 'Fuel Type', ylabel='Frequency of fuel type')

plt.subplot(1,3,3)
plt1 = car_dataset.carbody.value_counts().plot(kind='bar')
plt.title('Type Of Cars')
plt1.set(xlabel = 'Car Type', ylabel='Frequency of Car type')

plt.show()

```



```
car_dataset.head()
```

| car_ID | symboling | CompanyName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation |
|--------|-----------|-------------|-------------|------------|------------|---------|-------------|----------------|
| 0 | 1 | 3 | alfa-romero | gas | std | two | convertible | rwd |
| 1 | 2 | 3 | alfa-romero | gas | std | two | convertible | rwd |
| 2 | 3 | 1 | alfa-romero | gas | std | two | hatchback | rwd |
| 3 | 4 | 2 | audi | gas | std | four | sedan | fwd |
| 4 | 5 | 2 | audi | gas | std | four | sedan | 4wd |

```

plt.figure(figsize=(25,8))

plt.subplot(1,2,1)
plt.title('Symboling Histogram')

```

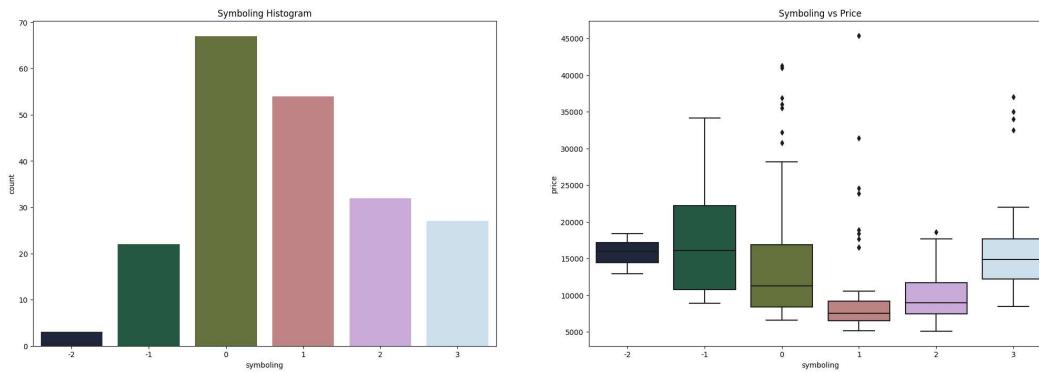
```

sns.countplot(car_dataset.symboling, x=car_dataset.symboling, palette=("cubehelix"))

plt.subplot(1,2,2)
plt.title('Symboling vs Price')
sns.boxplot(x=car_dataset.symboling, y=car_dataset.price, palette=("cubehelix"))

plt.show()

```



```

plt.figure(figsize=(25,8))

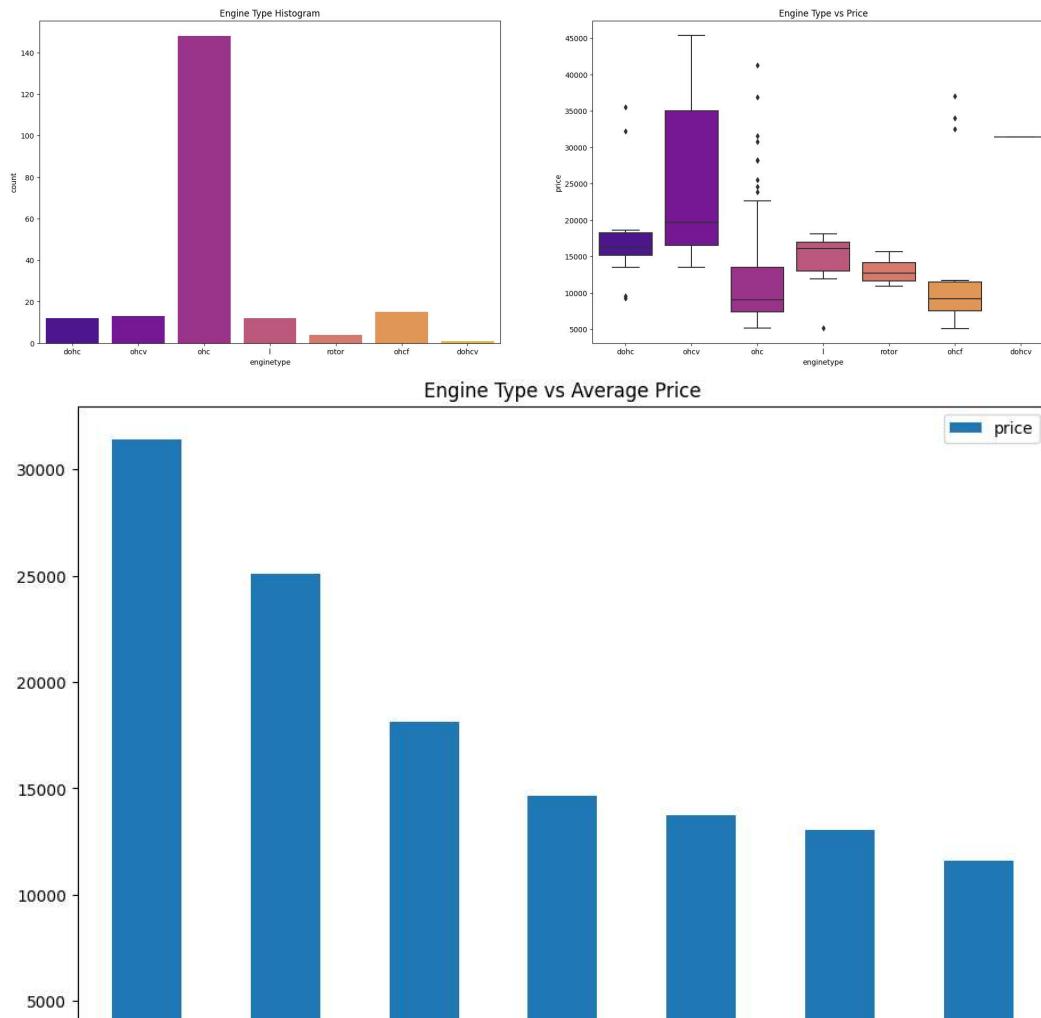
plt.subplot(1,2,1)
plt.title('Engine Type Histogram')
sns.countplot(car_dataset.enginetype, x=car_dataset.enginetype, palette=("plasma"))

plt.subplot(1,2,2)
plt.title('Engine Type vs Price')
sns.boxplot(x=car_dataset.enginetype, y=car_dataset.price, palette=("plasma"))

plt.show()

df = pd.DataFrame(car_dataset.groupby(['enginetype'])['price'].mean().sort_values(ascending = False))
df.plot.bar(figsize=(11,8))
plt.title('Engine Type vs Average Price')
plt.show()

```



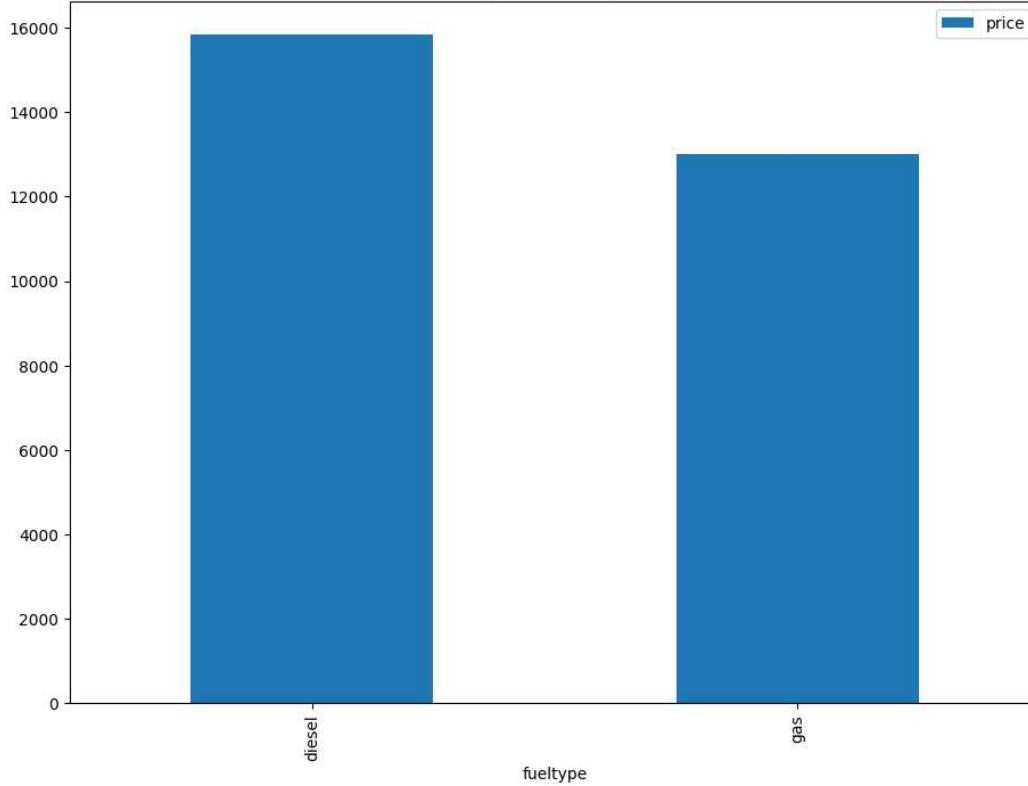
```
df = pd.DataFrame(car_dataset.groupby(['CompanyName'])['price'].mean().sort_values(ascending = False))
df.plot.bar(figsize=(11,8))
plt.title('Company Name vs Average Price')
plt.show()
```

Company Name vs Average Price



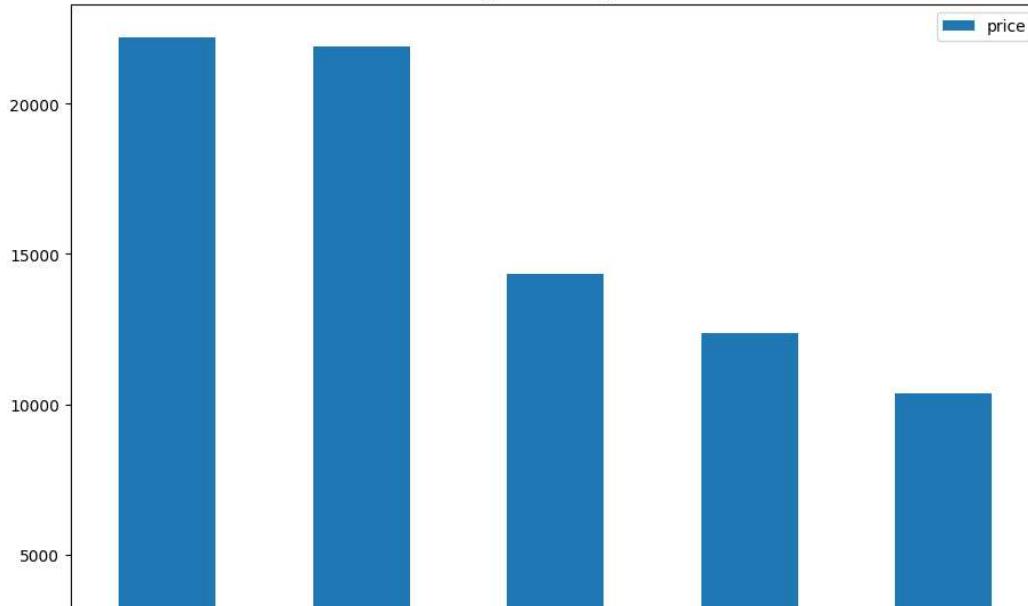
```
df = pd.DataFrame(car_dataset.groupby(['fueltype'])['price'].mean().sort_values(ascending = False))
df.plot.bar(figsize=(11,8))
plt.title('Fuel Type vs Average Price')
plt.show()
```

Fuel Type vs Average Price



```
df = pd.DataFrame(car_dataset.groupby(['carbody'])['price'].mean().sort_values(ascending = False))
df.plot.bar(figsize=(11,8))
plt.title('Car Type vs Average Price')
plt.show()
```

Car Type vs Average Price



```

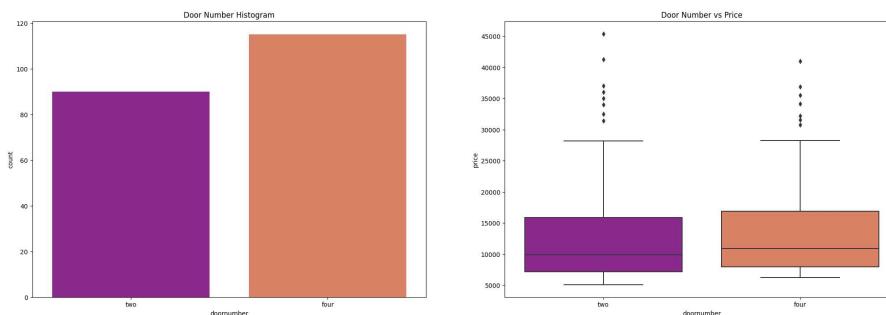
plt.figure(figsize=(25,8))

plt.subplot(1,2,1)
plt.title('Door Number Histogram')
sns.countplot(x=car_dataset.doornumber, palette=("plasma"))

plt.subplot(1,2,2)
plt.title('Door Number vs Price')
sns.boxplot(x=car_dataset.doornumber, y=car_dataset.price, palette=("plasma"))

plt.show()

```



```

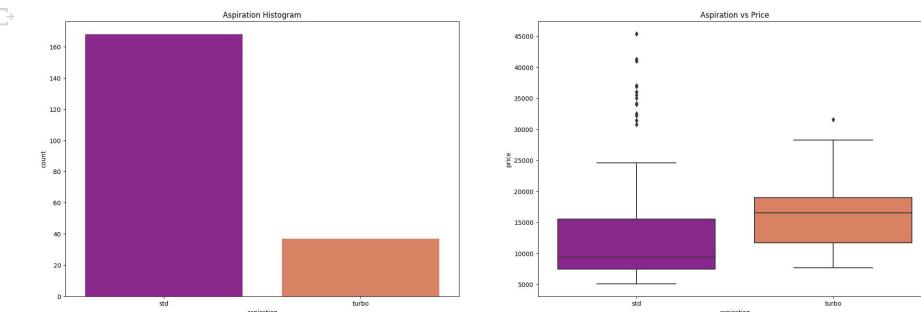
plt.figure(figsize=(25,8))

plt.subplot(1,2,1)
plt.title('Aspiration Histogram')
sns.countplot(x=car_dataset.aspiration, palette=("plasma"))

plt.subplot(1,2,2)
plt.title('Aspiration vs Price')
sns.boxplot(x=car_dataset.aspiration, y=car_dataset.price, palette=("plasma"))

plt.show()

```



```
np.corrcoef(car_dataset['carlength'], car_dataset['carwidth'])[0, 1]
0.841118268481846

#fueleconomy
car_dataset['fueleconomy'] = (0.55 * car_dataset['citympg']) + (0.45 * car_dataset['highwaympg'])

#Binning the Car Companies based on avg prices of each Company.
car_dataset['price'] = car_dataset['price'].astype('int')
temp = car_dataset.copy()
table = temp.groupby(['CompanyName'])['price'].mean()
temp = temp.merge(table.reset_index(), how='left', on='CompanyName')
bins = [0,10000,20000,40000]
cars_bin=['Budget','Medium','Highend']
car_dataset['carsrange'] = pd.cut(temp['price_y'],bins,right=False,labels=cars_bin)
```

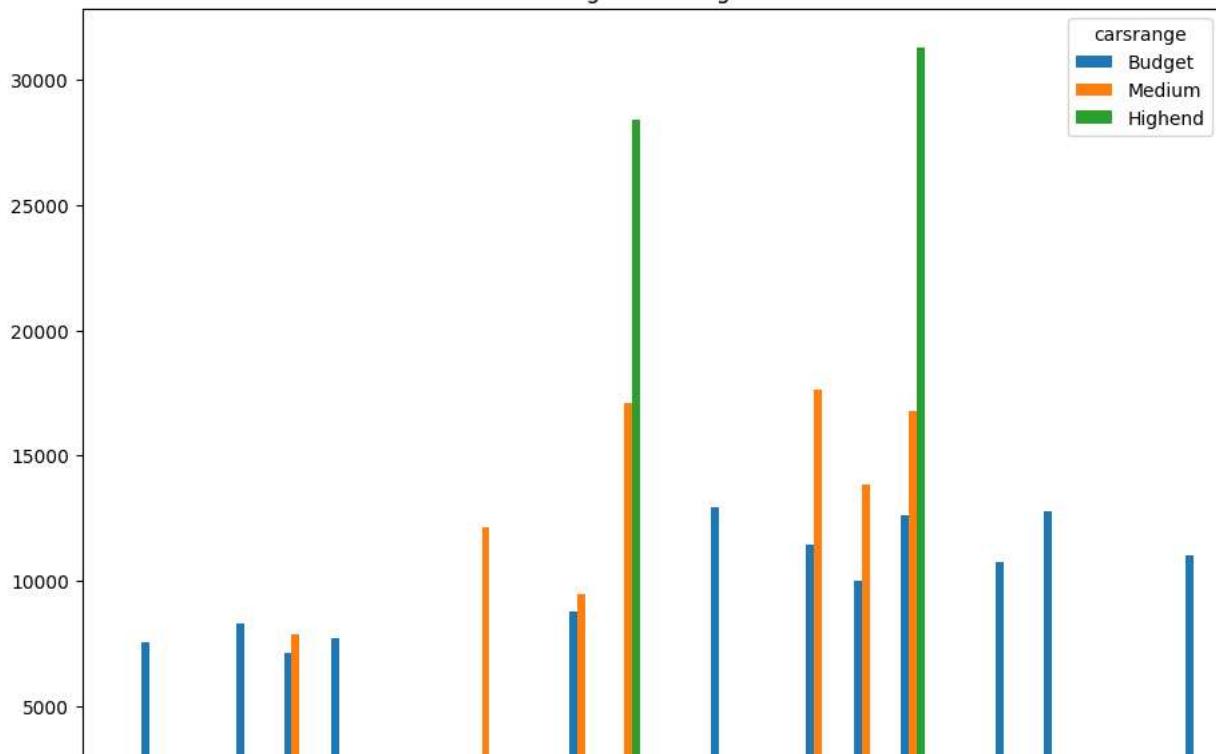
```
car_dataset.head()
```

| car_ID | symboling | CompanyName | fueltype | aspiration | doornumber | carbody | drivewheel | enginelocation | wheelbase | carlength | carwidt |
|--------|-----------|-------------|-------------|------------|------------|---------|-------------|----------------|-----------|-----------|---------|
| 0 | 1 | 3 | alfa-romero | gas | std | two | convertible | rwd | front | 88.6 | 168.8 |
| 1 | 2 | 3 | alfa-romero | gas | std | two | convertible | rwd | front | 88.6 | 168.8 |
| 2 | 3 | 1 | alfa-romero | gas | std | two | hatchback | rwd | front | 94.5 | 171.2 |
| 3 | 4 | 2 | audi | gas | std | four | sedan | fwd | front | 99.8 | 176.6 |
| 4 | 5 | 2 | audi | gas | std | four | sedan | 4wd | front | 99.4 | 176.6 |

```
df = pd.DataFrame(car_dataset.groupby(['fuelsystem','drivewheel','carsrange'])['price'].mean().unstack(fill_value=0))
df.plot.bar(figsize=(11,8))
plt.title('Car Range vs Average Price')

plt.show()
```

Car Range vs Average Price



```
cars_lr = car_dataset[['price', 'fueltype', 'aspiration', 'carbody', 'drivewheel', 'wheelbase',
                      'curbweight', 'enginetype', 'cylindernumber', 'enginesize', 'boreratio', 'horsepower',
                      'fueleconomy', 'carlength', 'carwidth', 'carsrange']]
cars_lr.head()
```

| | price | fueltype | aspiration | carbody | drivewheel | wheelbase | curbweight | enginetype | cylindernumber | enginesize | boreratio | horsepow |
|---|-------|----------|------------|-------------|------------|-----------|------------|------------|----------------|------------|-----------|----------|
| 0 | 13495 | gas | std | convertible | rwd | 88.6 | 2548 | dohc | four | 130 | 3.47 | 1 |
| 1 | 16500 | gas | std | convertible | rwd | 88.6 | 2548 | dohc | four | 130 | 3.47 | 1 |
| 2 | 16500 | gas | std | hatchback | rwd | 94.5 | 2823 | ohcv | six | 152 | 2.68 | 1 |
| 3 | 13950 | gas | std | sedan | fwd | 99.8 | 2337 | ohc | four | 109 | 3.19 | 1 |
| 4 | 17450 | gas | std | sedan | 4wd | 99.4 | 2824 | ohc | five | 136 | 3.19 | 1 |

```
# Defining the map function
def dummies(x,df):
    temp = pd.get_dummies(df[x], drop_first = True)
    df = pd.concat([df, temp], axis = 1)
    df.drop([x], axis = 1, inplace = True)
    return df
```

```
# Applying the function to the cars_lr
```

```
cars_lr = dummies('fueltype',cars_lr)
cars_lr = dummies('aspiration',cars_lr)
cars_lr = dummies('carbody',cars_lr)
cars_lr = dummies('drivewheel',cars_lr)
cars_lr = dummies('enginetype',cars_lr)
cars_lr = dummies('cylindernumber',cars_lr)
cars_lr = dummies('carsrange',cars_lr)
```

```
cars_lr.head()
```

| | price | wheelbase | curbweight | enginesize | boreratio | horsepower | fueleconomy | carlength | carwidth | gas | turbo | hardtop | hatchback | sed |
|---|-------|-----------|------------|------------|-----------|------------|-------------|-----------|----------|-----|-------|---------|-----------|-----|
| 0 | 13495 | 88.6 | 2548 | 130 | 3.47 | 111 | 23.70 | 168.8 | 64.1 | 1 | 0 | 0 | 0 | |
| 1 | 16500 | 88.6 | 2548 | 130 | 3.47 | 111 | 23.70 | 168.8 | 64.1 | 1 | 0 | 0 | 0 | |

```
cars_lr.shape
```

```
(205, 31)
```

```
#Train-Test
from sklearn.model_selection import train_test_split

np.random.seed(0)
df_train, df_test = train_test_split(cars_lr, train_size = 0.7, test_size = 0.3, random_state = 100)

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
num_vars = ['wheelbase', 'curbweight', 'enginesize', 'boreratio', 'horsepower','fueleconomy','carlength','carwidth','price']
df_train[num_vars] = scaler.fit_transform(df_train[num_vars])

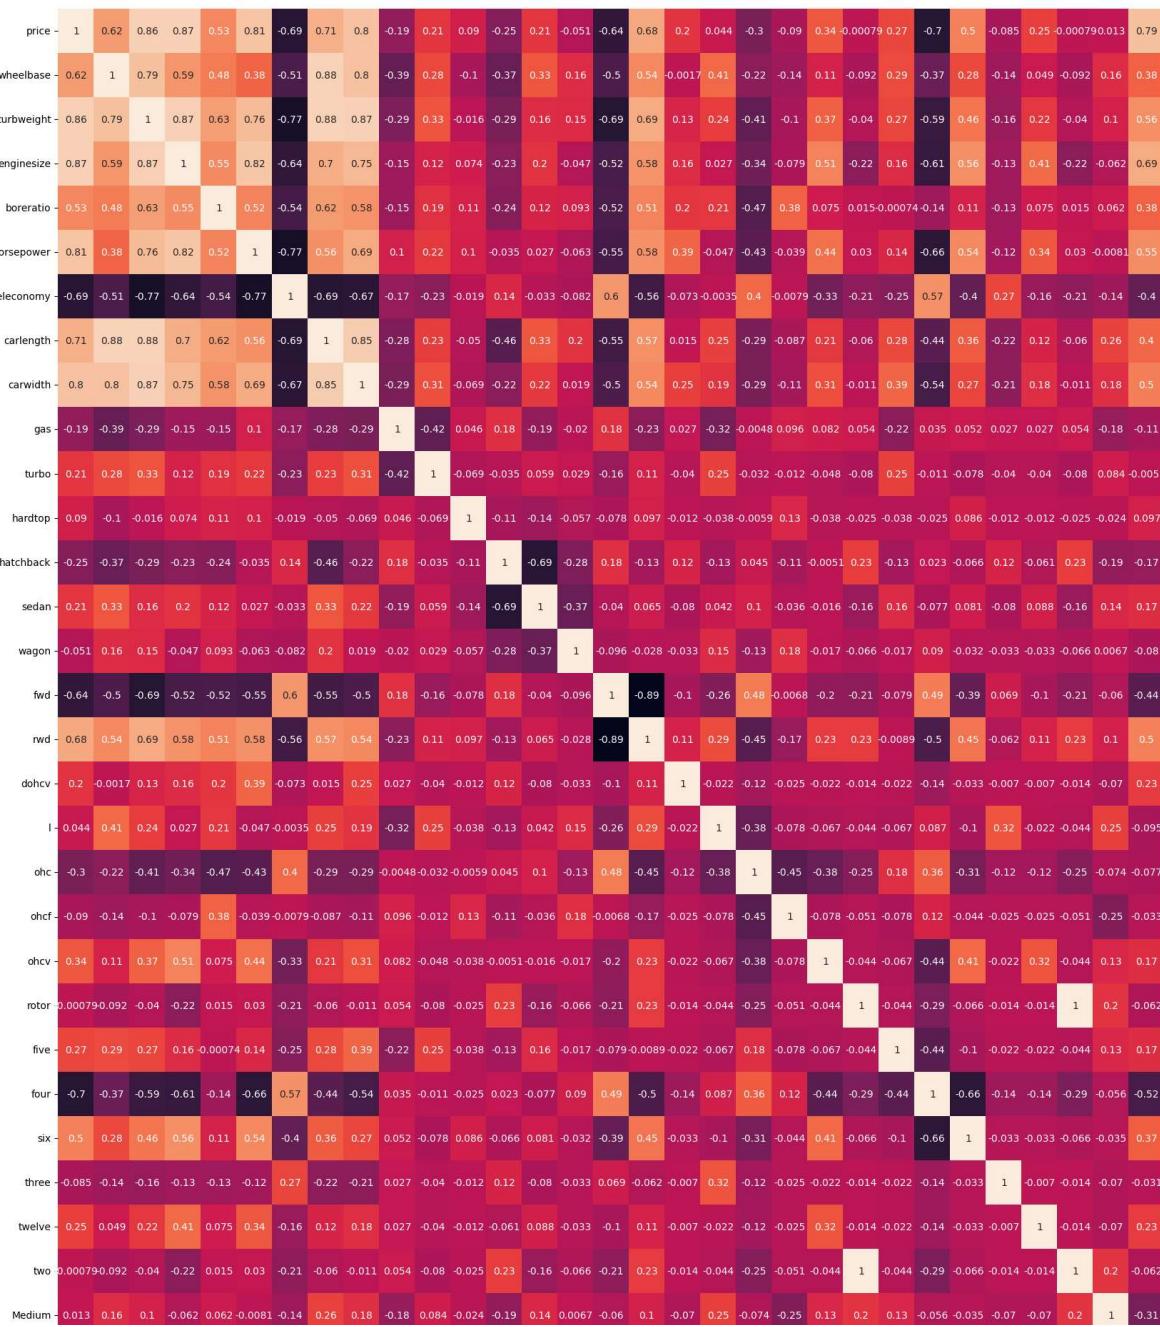
df_train.head()
```

| | price | wheelbase | curbweight | enginesize | boreratio | horsepower | fueleconomy | carlength | carwidth |
|-----|----------|-----------|------------|------------|-----------|------------|-------------|-----------|----------|
| 122 | 0.068818 | 0.244828 | 0.272692 | 0.139623 | 0.230159 | 0.083333 | 0.530864 | 0.426016 | 0.291666 |
| 125 | 0.466890 | 0.272414 | 0.500388 | 0.339623 | 1.000000 | 0.395833 | 0.213992 | 0.452033 | 0.666666 |
| 166 | 0.122110 | 0.272414 | 0.314973 | 0.139623 | 0.444444 | 0.266667 | 0.344307 | 0.448780 | 0.308333 |
| 1 | 0.314446 | 0.068966 | 0.411171 | 0.260377 | 0.626984 | 0.262500 | 0.244170 | 0.450407 | 0.316666 |
| 199 | 0.382131 | 0.610345 | 0.647401 | 0.260377 | 0.746032 | 0.475000 | 0.122085 | 0.775610 | 0.575000 |

```
df_train.describe()
```

| | price | wheelbase | curbweight | enginesize | boreratio | horsepower | fueleconomy | carlength | carwidth | gas | turbo |
|-------|------------|------------|------------|------------|------------|------------|-------------|------------|------------|------------|------------|
| count | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 | 143.000000 |
| mean | 0.219309 | 0.411141 | 0.407878 | 0.241351 | 0.497946 | 0.227302 | 0.358265 | 0.525476 | 0.461655 | 0.909091 | 0.181818 |
| std | 0.215682 | 0.205581 | 0.211269 | 0.154619 | 0.207140 | 0.165511 | 0.185980 | 0.204848 | 0.184517 | 0.288490 | 0.387050 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.067298 | 0.272414 | 0.245539 | 0.135849 | 0.305556 | 0.091667 | 0.198903 | 0.399187 | 0.304167 | 1.000000 | 0.000000 |
| 50% | 0.140343 | 0.341379 | 0.355702 | 0.184906 | 0.500000 | 0.191667 | 0.344307 | 0.502439 | 0.425000 | 1.000000 | 0.000000 |
| 75% | 0.313479 | 0.503448 | 0.559542 | 0.301887 | 0.682540 | 0.283333 | 0.512346 | 0.669919 | 0.550000 | 1.000000 | 0.000000 |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |

```
plt.figure(figsize = (25, 25))
sns.heatmap(df_train.corr(), annot = True)
plt.show()
```



```
#Model Building
#Dividing data into X and y variables
y_train = df_train.pop('price')
X_train = df_train
```

```
lm = LinearRegression()
lm.fit(X_train,y_train)
rfe = RFE(lm)
```

```
rfe = rfe.fit(X_train, y_train)
```

```
list(zip(X_train.columns,rfe.support_,rfe.ranking_))
```

```
[('wheelbase', True, 1),
 ('curbweight', True, 1),
 ('enginesize', False, 8),
 ('boreratio', False, 5),
 ('horsepower', True, 1),
 ('fueleconomy', True, 1),
 ('carlength', False, 6),
 ('carwidth', True, 1),
 ('gas', False, 12),
 ('turbo', False, 13),
 ('hardtop', True, 1),
 ('hatchback', True, 1),
 ('sedan', True, 1),
 ('wagon', True, 1),
 ('fwd', False, 11),
 ('rwd', False, 10),
 ('dohcv', True, 1),
 ('l', False, 14),
 ('ohc', False, 2),
 ('ohcf', False, 3),
 ('ohcv', False, 4),
 ('rotor', False, 16),
 ('five', True, 1),
 ('four', True, 1),
 ('six', True, 1),
 ('three', False, 9),
 ('twelve', True, 1),
 ('two', False, 15),
 ('Medium', False, 7),
 ('Highend', True, 1)]
```

```
X_train.columns[rfe.support_]
```

```
Index(['wheelbase', 'curbweight', 'horsepower', 'fueleconomy', 'carwidth',
       'hardtop', 'hatchback', 'sedan', 'wagon', 'dohcv', 'five', 'four',
       'six', 'twelve', 'Highend'],
      dtype='object')
```

```
X_train_rfe = X_train[X_train.columns[rfe.support_]]
```

```
X_train_rfe.head()
```

| | wheelbase | curbweight | horsepower | fueleconomy | carwidth | hardtop | hatchback | sedan | wagon | dohcv | five | four | six | twelve | Highend |
|-----|-----------|------------|------------|-------------|----------|---------|-----------|-------|-------|-------|------|------|-----|--------|---------|
| 122 | 0.244828 | 0.272692 | 0.083333 | 0.530864 | 0.291667 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| 125 | 0.272414 | 0.500388 | 0.395833 | 0.213992 | 0.666667 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| 166 | 0.272414 | 0.314973 | 0.266667 | 0.344307 | 0.308333 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| 1 | 0.068966 | 0.411171 | 0.262500 | 0.244170 | 0.316667 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| 199 | 0.610345 | 0.647401 | 0.475000 | 0.122085 | 0.575000 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | 0 |

```
def build_model(X,y):
    X = sm.add_constant(X)
    lm = sm.OLS(y,X).fit()
    print(lm.summary())
    return X
```

```
def checkVIF(X):
    vif = pd.DataFrame()
    vif['Features'] = X.columns
    vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
    vif['VIF'] = round(vif['VIF'], 2)
    vif = vif.sort_values(by = "VIF", ascending = False)
    return(vif)
```

```
X_train_new = build_model(X_train_rfe,y_train)
```

```
OLS Regression Results
=====
Dep. Variable:          price    R-squared:           0.936
Model:                 OLS     Adj. R-squared:        0.929
Method:                Least Squares F-statistic:         124.4
Date: Fri, 11 Aug 2023 Prob (F-statistic):      3.16e-68
Time: 14:00:32 Log-Likelihood:            213.80
No. Observations:      143     AIC:                  -395.6
Df Residuals:          127     BIC:                  -348.2
Df Model:               15
Covariance Type:       nonrobust
=====
      coef    std err     t   P>|t|    [0.025    0.975]
-----
const    0.0108   0.050   0.214   0.831   -0.089   0.110
wheelbase 0.0814   0.063   1.298   0.197   -0.043   0.205
curbweight 0.2017   0.081   2.482   0.014   0.041   0.362
horsepower 0.4801   0.098   4.923   0.000   0.287   0.673
fueleconomy 0.1300   0.054   2.418   0.017   0.024   0.236
carwidth  0.2140   0.074   2.903   0.004   0.068   0.360
hardtop   -0.0611   0.048  -1.278   0.204   -0.156   0.034
hatchback -0.1379   0.035  -3.977   0.000   -0.206   -0.069
sedan     -0.1119   0.036  -3.133   0.002   -0.183   -0.041
wagon     -0.1350   0.036  -3.708   0.000   -0.207   -0.063
dohcv     -0.2880   0.086  -3.360   0.001   -0.458   -0.118
five      -0.0416   0.031  -1.346   0.181   -0.103   0.020
four      -0.0802   0.025  -3.249   0.001   -0.129   -0.031
six       -0.0425   0.030  -1.400   0.164   -0.103   0.018
twelve    -0.1348   0.075  -1.790   0.076   -0.284   0.014
Highend   0.2416   0.020  12.095   0.000   0.202   0.281
=====
Omnibus:             54.936 Durbin-Watson:          2.013
Prob(Omnibus):        0.000 Jarque-Bera (JB):      231.186
Skew:                 1.331 Prob(JB):            6.29e-51
Kurtosis:            8.632 Cond. No.             46.8
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
X_train_new = X_train_new.drop(["fueleconomy"], axis = 1)
```

```
X_train_new = build_model(X_train_new,y_train)
```

```
OLS Regression Results
=====
Dep. Variable:          price    R-squared:           0.933
Model:                 OLS     Adj. R-squared:        0.926
Method:                Least Squares F-statistic:         128.0
Date: Fri, 11 Aug 2023 Prob (F-statistic):      4.67e-68
Time: 14:00:53 Log-Likelihood:            210.58
No. Observations:      143     AIC:                  -391.2
Df Residuals:          128     BIC:                  -346.7
Df Model:               14
Covariance Type:       nonrobust
=====
      coef    std err     t   P>|t|    [0.025    0.975]
-----
const    0.0739   0.044   1.690   0.093   -0.013   0.160
wheelbase 0.0594   0.063   0.940   0.349   -0.066   0.184
curbweight 0.1875   0.083   2.271   0.025   0.024   0.351
horsepower 0.3425   0.081   4.246   0.000   0.183   0.502
carwidth  0.2264   0.075   3.023   0.003   0.078   0.375
hardtop   -0.0430   0.048  -0.895   0.373   -0.138   0.052
hatchback -0.1256   0.035  -3.596   0.000   -0.195   -0.056
sedan     -0.1001   0.036  -2.778   0.006   -0.171   -0.029
wagon     -0.1288   0.037  -3.482   0.001   -0.202   -0.056
dohcv     -0.2002   0.079  -2.532   0.013   -0.357   -0.044
five      -0.0378   0.031  -1.203   0.231   -0.100   0.024
four      -0.0694   0.025  -2.807   0.006   -0.118   -0.020
six       -0.0216   0.030  -0.729   0.467   -0.080   0.037
twelve    -0.0793   0.073  -1.085   0.280   -0.224   0.065
Highend   0.2484   0.020  12.327   0.000   0.209   0.288
=====
Omnibus:             53.309 Durbin-Watson:          2.045
Prob(Omnibus):        0.000 Jarque-Bera (JB):      212.461
Skew:                 1.307 Prob(JB):            7.32e-47
Kurtosis:            8.369 Cond. No.             39.7
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

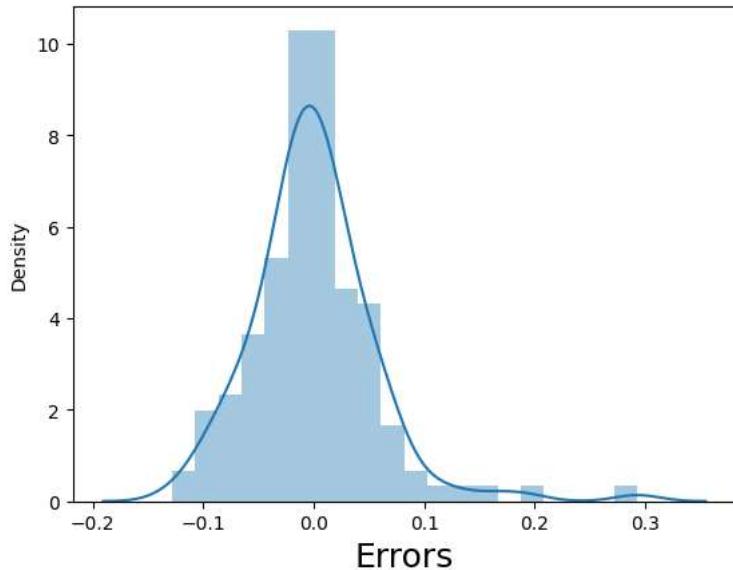
lm = sm.OLS(y_train,X_train_new).fit()
y_train_price = lm.predict(X_train_new)

fig = plt.figure()
sns.distplot((y_train - y_train_price), bins = 20)
fig.suptitle('Error Terms', fontsize = 20)
plt.xlabel('Errors', fontsize = 18)

Text(0.5, 0, 'Errors')

```

Error Terms



```
print(lm.summary())
```

```

OLS Regression Results
=====
Dep. Variable:      price   R-squared:       0.933
Model:              OLS    Adj. R-squared:     0.926
Method:             Least Squares F-statistic:     128.0
Date: Fri, 11 Aug 2023 Prob (F-statistic): 4.67e-68
Time: 14:01:22 Log-Likelihood: 210.58
No. Observations: 143    AIC:            -391.2
Df Residuals:     128    BIC:            -346.7
Df Model:         14
Covariance Type:  nonrobust
=====
            coef    std err        t      P>|t|      [0.025      0.975]
-----
const    0.0739   0.044     1.690     0.093     -0.013     0.160
wheelbase 0.0594   0.063     0.940     0.349     -0.066     0.184
curbweight 0.1875   0.083     2.271     0.025     0.024     0.351
horsepower 0.3425   0.081     4.246     0.000     0.183     0.502
carwidth  0.2264   0.075     3.023     0.003     0.078     0.375
hardtop   -0.0430  0.048     -0.895    0.373     -0.138     0.052
hatchback -0.1256  0.035     -3.596    0.000     -0.195    -0.056
sedan    -0.1001  0.036     -2.778    0.006     -0.171    -0.029
wagon    -0.1288  0.037     -3.482    0.001     -0.202    -0.056
dohcv    -0.2002  0.079     -2.532    0.013     -0.357    -0.044
five     -0.0378  0.031     -1.203    0.231     -0.100     0.024
four     -0.0694  0.025     -2.807    0.006     -0.118    -0.020
six      -0.0216  0.030     -0.729    0.467     -0.080     0.037
twelve   -0.0793  0.073     -1.085    0.280     -0.224     0.065
Highend  0.2484   0.020     12.327    0.000     0.209     0.288
=====
Omnibus:           53.309 Durbin-Watson:      2.045
Prob(Omnibus):    0.000  Jarque-Bera (JB): 212.461
Skew:               1.307 Prob(JB):        7.32e-47
Kurtosis:          8.369 Cond. No.          39.7
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.